

Simplicio-Loop vs Baseline — Issue #17 Benchmark

Task: P0 slice of simplicio-agent#17 (Asolaria HRM + N-Nest-Prime): Brown-Hilbert addressing + Watcher corrective gate.

Environment: loop=3.38.9 mapper=0.26.0 cli=0.18.0 fast=2.0.17

Metric	Without loop	With loop	With/Without
Wall time (s)	2.72	13.20	4.857
Est. tokens (chars/4)	156576	15246	0.097
Context bytes	626303	60982	0.097
Commands	16	12	
Operator calls	0	6	
Success	True	True	
Token savings %	—	90.3	(positive = loop uses fewer est. tokens)

Phase times (seconds)

Phase	Without	With
implement	0.00	0.00
plan	0.00	1.99
preflight	0.00	2.78
survey	1.50	7.31
test	1.21	1.13

Method notes

Without loop: ad-hoc rg + full file reads (LLM-style survey). No preflight/mapper/fast.

With loop: STRICT env, preflight --strict, mapper scan/inspect/handoff (decide context capped), fast doctor, arm_drain_prism, signature-only reads.

Deliverable (both paths): agent/asolaria_nest_contract.py + tests (NestAddress + watcher_gate). Implementation is identical so the comparison isolates orchestration/survey cost. Token figures are **estimated** as context_bytes/4 (not provider-billed tokens).

Interpretation

Loop path reduced estimated intake tokens by ~90.3% (156576 → 15246). Wall time may be higher due to operator startup (preflight/mapper) on a cold run; token economy is the primary win for multi-turn LLM sessions.

Artifacts: C:/Users/Z0059V7A/AppData/Local/Temp/simplicio-loop-bench-issue-17 · Generated
UTC wall clock embedded in metrics JSON.